


```
from google.colab import files
uploaded=files.upload()
```

 Choose Files seeds.csv

- **seeds.csv**(text/csv) - 9498 bytes, last modified: 7/7/2025 - 100% done

Saving seeds.csv to seeds (1).csv

```
import pandas
import numpy
import matplotlib.pyplot as plt
```

```
from sklearn import model_selection
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import StandardScaler
```


```
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

```
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
```


```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
```

```
dataset = pandas.read_csv("/content/seeds.csv")
print(dataset.head())
```



	area	parameter	compact	length	width	asymmetry	groove	class
0	15.26	14.84	0.871	5.763	3.312	2.221	5.22	1
1	14.88	14.57	0.8811	5.554	3.333	1.018	4.956	1
2	14.29	14.09	0.9050	5.291	3.337	2.699	4.825	1
3	13.84	13.94	0.8955	5.324	3.379	2.259	4.805	1
4	16.14	14.99	0.9034	5.658	3.562	1.355	5.175	1


```
# load the data assign names
url="seeds.csv"
names=['area','parameter','compact','length','width','asymmetry','groove','class']
dataset=pandas.read_csv(url,names=names)
#Data
print(dataset)
```



	area	parameter	compact	length	width	asymmetry	groove	class
0	15.26	14.84	0.8710	5.763	3.312	2.221	5.220	1
1	14.88	14.57	0.8811	5.554	3.333	1.018	4.956	1
2	14.29	14.09	0.9050	5.291	3.337	2.699	4.825	1
3	13.84	13.94	0.8955	5.324	3.379	2.259	4.805	1
4	16.14	14.99	0.9034	5.658	3.562	1.355	5.175	1
..
205	12.19	13.20	0.8783	5.137	2.981	3.631	4.870	3
206	11.23	12.88	0.8511	5.140	2.795	4.325	5.003	3
207	13.20	13.66	0.8883	5.236	3.232	8.315	5.056	3
208	11.84	13.21	0.8521	5.175	2.836	3.598	5.044	3
209	12.30	13.34	0.8684	5.243	2.974	5.637	5.063	3

[210 rows x 8 columns]

```
# drescriptive statsfor data
print(dataset.dtypes)
```



	area	parameter	compact
	float64	float64	float64

```

length      float64
width       float64
asymmetry   float64
groove      float64
class       int64
dtype: object

```

```

print("data shepe",dataset.shape)
print("head of dataset",dataset.head(3))
print("tail of dataset",dataset.tail(3))

```

```

↗ data shepe (210, 8)
head of dataset      area  parameter  compact  length  width  asymmetry  groove  class
0  15.26      14.84    0.8710    5.763  3.312    2.221  5.220      1
1  14.88      14.57    0.8811    5.554  3.333    1.018  4.956      1
2  14.29      14.09    0.9050    5.291  3.337    2.699  4.825      1
tail of dataset      area  parameter  compact  length  width  asymmetry  groove  class
207 13.20      13.66    0.8883    5.236  3.232    8.315  5.056      3
208 11.84      13.21    0.8521    5.175  2.836    3.598  5.044      3
209 12.30      13.34    0.8684    5.243  2.974    5.637  5.063      3

```

```
dataset.apply(lambda x:sum(x.isnull()),axis=0)
```

```

↗
      0
area    0
parameter  0
compact  0
length    0
width     0
asymmetry 0
groove    0
class     0

```

dtype: int64

```

pandas.set_option('display.precision',2)
print(dataset.describe)

```

```

↗ <bound method NDFrame.describe of
0   15.26    14.84    0.87    5.76    3.31    2.22    5.22    1
1   14.88    14.57    0.88    5.55    3.33    1.02    4.96    1
2   14.29    14.09    0.91    5.29    3.34    2.70    4.83    1
3   13.84    13.94    0.90    5.32    3.38    2.26    4.80    1
4   16.14    14.99    0.90    5.66    3.56    1.35    5.17    1
..    ...    ...    ...    ...    ...    ...    ...    ...
205 12.19    13.20    0.88    5.14    2.98    3.63    4.87    3
206 11.23    12.88    0.85    5.14    2.79    4.33    5.00    3
207 13.20    13.66    0.89    5.24    3.23    8.31    5.06    3
208 11.84    13.21    0.85    5.17    2.84    3.60    5.04    3
209 12.30    13.34    0.87    5.24    2.97    5.64    5.06    3


```

[210 rows x 8 columns]>

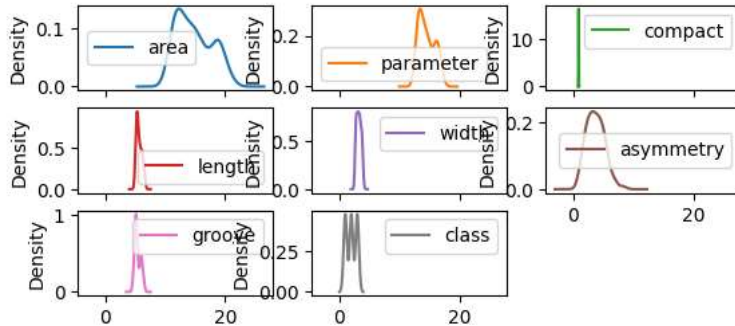
```

dataset.hist( )
plt.show()

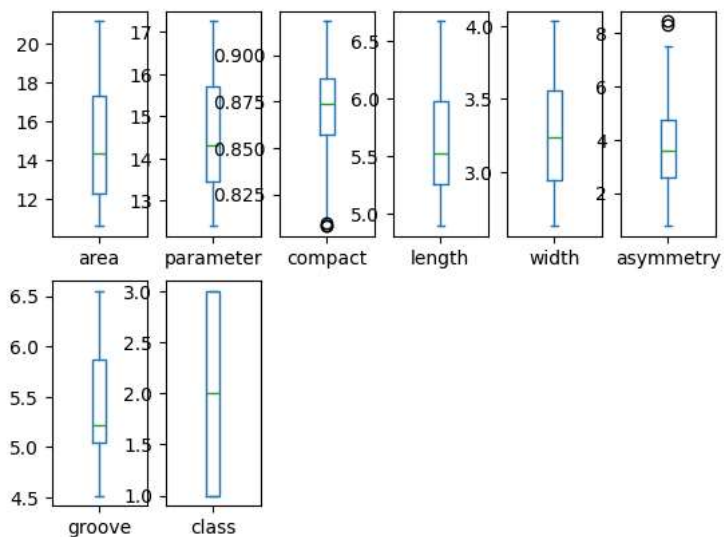
```



```
dataset.plot(kind='density',subplots=True,layout=(5,3))
plt.show()
```



```
#to find outliers in dataset's attributes
dataset.plot(kind='box',subplots=True,layout=(2,6))
plt.show()
```




```
scatter_matrix(dataset)
plt.show()
```



```
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
```

```
results = []
names = []
scoring = 'accuracy'
test_size = 0.33
seed = 8
```

```
for name, model in models:
    kfold = model_selection.ShuffleSplit(n_splits=10, test_size=test_size, random_state=seed)
    X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y, test_size=test_size, random_state=seed)
    cv_results = model_selection.cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean()*100, cv_results.std())
    print(msg)
```

 /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

```
fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```



Algorithm Comparison



```
pipelines = []
pipelines.append(('ScaledLR', Pipeline([('Scaler', StandardScaler()), ('LR', LogisticRegression())]))
pipelines.append(('ScaledLDA', Pipeline([('Scaler', StandardScaler()), ('LDA', LinearDiscriminantAnalysis())]))
pipelines.append(('ScaledKNN', Pipeline([('Scaler', StandardScaler()), ('KNN', KNeighborsClassifier())]))
pipelines.append(('ScaledCART', Pipeline([('Scaler', StandardScaler()), ('CART', DecisionTreeClassifier())]))
pipelines.append(('ScaledNB', Pipeline([('Scaler', StandardScaler()), ('NB', GaussianNB())]))
pipelines.append(('ScaledSVM', Pipeline([('Scaler', StandardScaler()), ('SVM', SVC())]))
results = []
names = []
for name, model in pipelines:
    kfold = model_selection.ShuffleSplit(n_splits=10, test_size=test_size, random_state=seed)
    #X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y, test_size=test_size, random_state=seed)
    cv_results = model_selection.cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean()*100, cv_results.std())
    print(msg)
```



```
ScaledLR: 93.285714 (0.015714)
ScaledLDA: 96.428571 (0.017202)
ScaledKNN: 92.571429 (0.029137)
ScaledCART: 89.428571 (0.037362)
ScaledNB: 90.285714 (0.026186)
ScaledSVM: 93.000000 (0.016225)
```

```
scaler = StandardScaler().fit(X) # KNN
rescaledX = scaler.transform(X)
neighbors = [1,3,5,7,9,11,13,15,17,19,21]
param_grid = dict(n_neighbors=neighbors)
model = KNeighborsClassifier()
kfold = model_selection.ShuffleSplit(n_splits=10, test_size=test_size, random_state=seed)
#X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y, test_size=test_size, random_state=seed)
grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring=scoring, cv=kfold)
grid_result = grid.fit(rescaledX, Y)
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```



```
Best: 0.931429 using {'n_neighbors': 1}
```

```
scaler = StandardScaler().fit(X) #SVM
rescaledX = scaler.transform(X)
c_values = [0.1, 0.3, 0.5, 0.7, 0.9, 1.0, 1.3, 1.5, 1.7, 2.0]
kernel_values = ['linear', 'poly', 'rbf', 'sigmoid']
```

```

param_grid = dict(C=c_values, kernel=kernel_values)
model = SVC()
kfold = model_selection.ShuffleSplit(n_splits=10, test_size=test_size, random_state=seed)
#X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y, test_size=test_size, random_state=seed)
grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring=scoring, cv=kfold)
grid_result = grid.fit(rescaledX, Y)
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))

```

Best: 0.941429 using {'C': 1.7, 'kernel': 'linear'}

```

ensembles = []
ensembles.append(('AB', AdaBoostClassifier()))
ensembles.append(('GBM', GradientBoostingClassifier()))
ensembles.append(('RF', RandomForestClassifier()))
ensembles.append(('ET', ExtraTreesClassifier()))
results = []
names = []
for name, model in ensembles:
    kfold = model_selection.ShuffleSplit(n_splits=10, test_size=test_size, random_state=seed)
    #X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y, test_size=test_size, random_state=seed)
    cv_results = model_selection.cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = " mean %s: %f ( standard deviation %f)" % (name, cv_results.mean()*100, cv_results.std())
    print(msg)

```


mean AB: 89.857143 (standard deviation 0.026688)
 mean GBM: 90.714286 (standard deviation 0.027293)
 mean RF: 91.000000 (standard deviation 0.028607)
 mean ET: 92.285714 (standard deviation 0.027255)

```

fig = plt.figure()
fig.suptitle('Ensemble Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```

Ensemble Algorithm Comparison



```

scaler = StandardScaler().fit(X);
rescaledX = scaler.transform(X)
model = LinearDiscriminantAnalysis()
model.fit(rescaledX, Y)

```

LinearDiscriminantAnalysis

```

rescaledValidationX = scaler.transform(X)
predictions = model.predict(rescaledValidationX)

```

```
print(accuracy_score(Y, predictions))
print(" ")
print(confusion_matrix(Y, predictions))
print(" ")
print(classification_report(Y,predictions))
```

0.9666666666666667

```
[[66  1  3]
 [ 0 70  0]
 [ 3  0 67]]
```

	precision	recall	f1-score	support
1.0	0.96	0.94	0.95	70
2.0	0.99	1.00	0.99	70
3.0	0.96	0.96	0.96	70
accuracy			0.97	210
macro avg	0.97	0.97	0.97	210
weighted avg	0.97	0.97	0.97	210